What's new in Cloud Tiering

Cloud Tiering

Ben Cammett, Aksel Davis
October 07, 2020

This PDF was generated from https://docs.netapp.com/us-en/cloud-tiering/reference_new.html on March 06, 2021. Always check docs.netapp.com for the latest.

Table of Contents

What's new in Cloud Tiering . . . 1
7 Sept 2020 . . . 1
5 Aug 2020 . . . 1
8 July 2020 . . . 2
8 June 2020 . . . 4
18 May 2020 . . . 5
1 Apr 2020 . . . 5
25 Dec 2019 . . . 6
3 Nov 2019 . . . 6
8 Sept 2019 . . . 7
7 Aug 2019 . . . 8
4 July 2019 . . . 8
10 June 2019 . . . 8
5 May 2019 . . . 9
7 Apr 2019 . . . 10

What's new in Cloud Tiering

NetApp periodically updates Cloud Tiering to bring you new features, enhancements, and bug fixes.

7 Sept 2020

This Cloud Tiering update includes the following enhancements.

Activate licenses (BYOL)

You can now activate a FabricPool license on an ONTAP cluster directly from Cloud Tiering. Learn more.

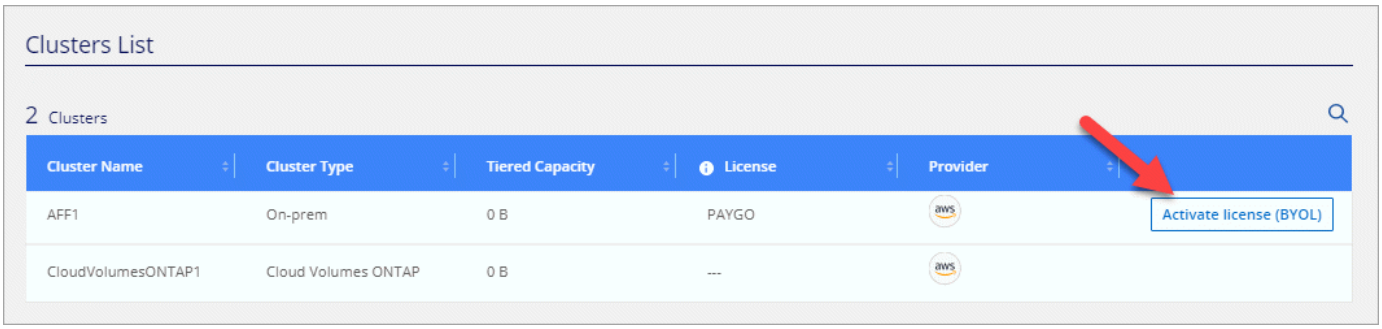

Azure resource group selection

A new step is available when you set up data tiering to Azure Blob storage. You now select a resource group first and then choose a Blob container.

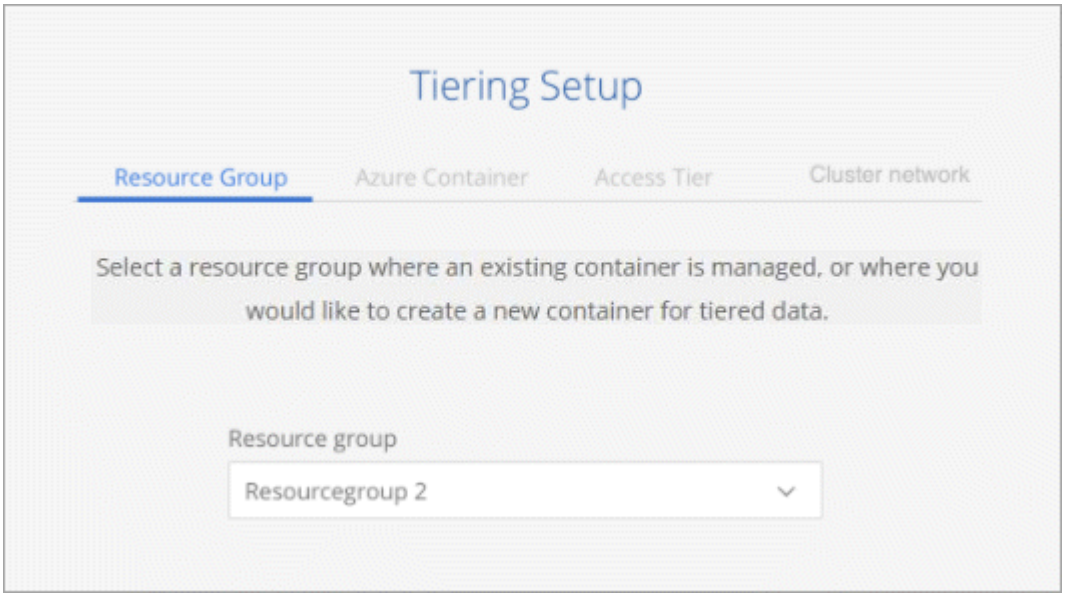

5 Aug 2020

If you have Cloud Volumes ONTAP systems, you'll find them in the Cluster Dashboard so you see a full view of data tiering in your hybrid cloud infrastructure.

From the Cluster Dashboard, you can view tiering information similar to an on-prem ONTAP cluster: operational health, current savings, savings opportunities, details about volumes and aggregates, and more.

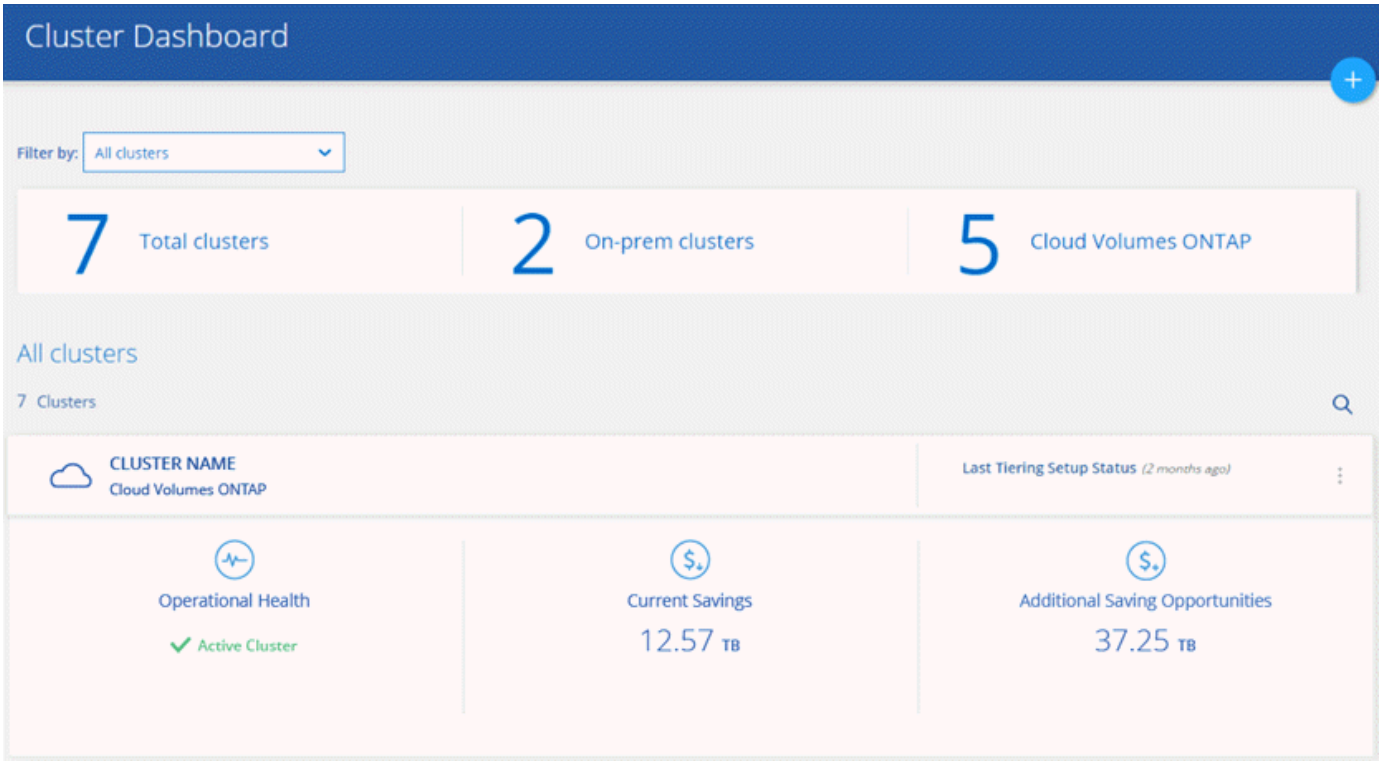

Cloud Volumes ONTAP systems are read-only from Cloud Tiering. You can't set up data tiering on Cloud Volumes ONTAP from Cloud Tiering. You'll still set up tiering the same way: from the working environment in Cloud Manager.

8 July 2020

This Cloud Tiering update includes the following new features and enhancements.

Tiering settings for aggregates

Each aggregate has two tiering settings that you can now adjust from Cloud Tiering: the tiering fullness threshold and whether inactive data reporting is enabled. Learn more.

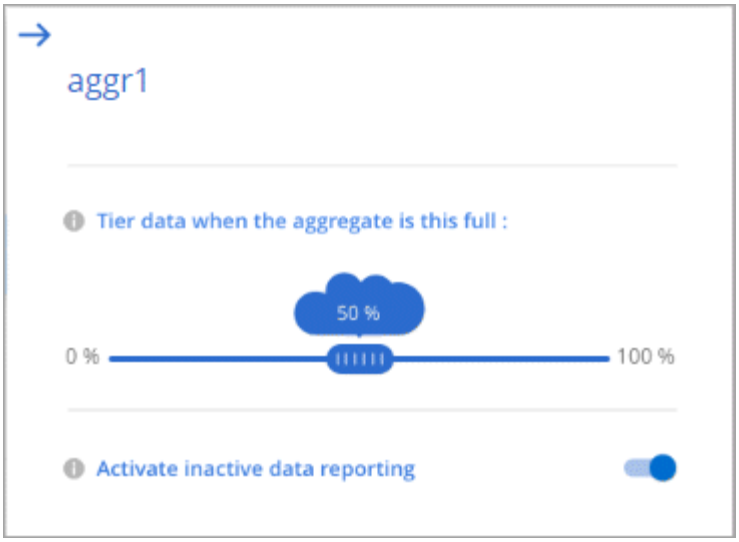

Inactive data reporting for discovered clusters

Inactive data reporting is now automatically enabled on ONTAP clusters that you discover through Cloud Manager. This enhancement makes it easier for Cloud Tiering to show you the potential savings from tiering cold data.

Remove clusters

If the health of a cluster is failed, you can now remove it from the dashboard.

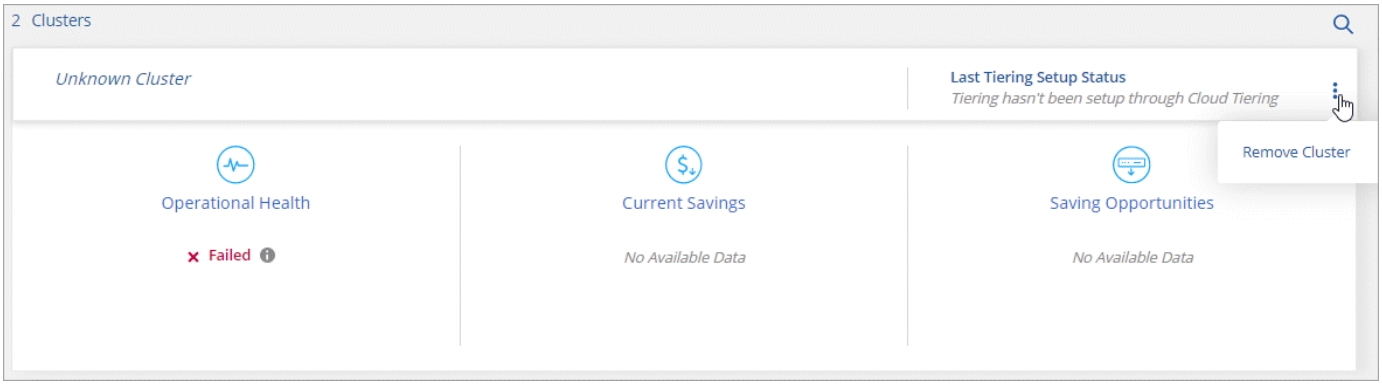

Discovery of unsupported clusters

Cloud Tiering now provides details about discovered clusters that aren't eligible for tiering.

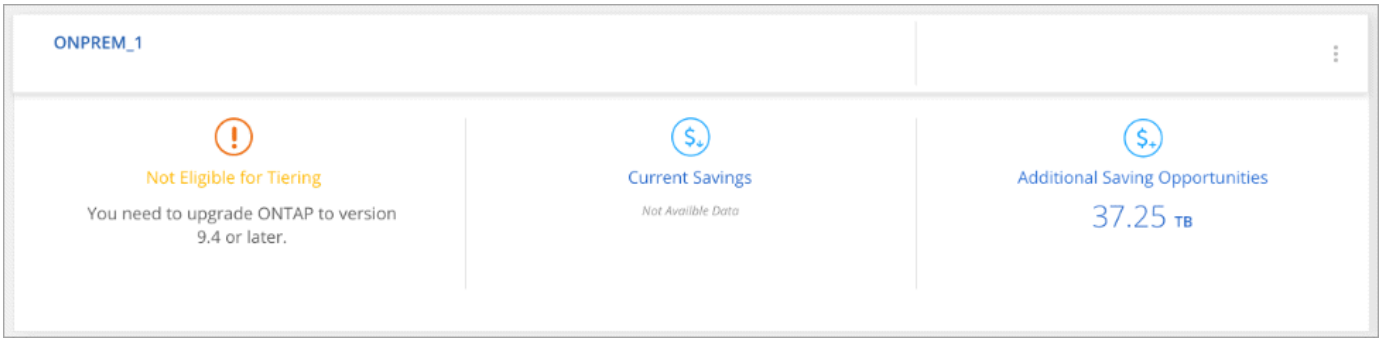

8 June 2020

This Cloud Tiering update includes the following new features and enhancements.

Support for StorageGRID

You can now tier inactive data from your ONTAP clusters to StorageGRID.

Learn how to get started.

View volumes before setup

You can now view details about a cluster's volumes before you set up tiering on the volumes.

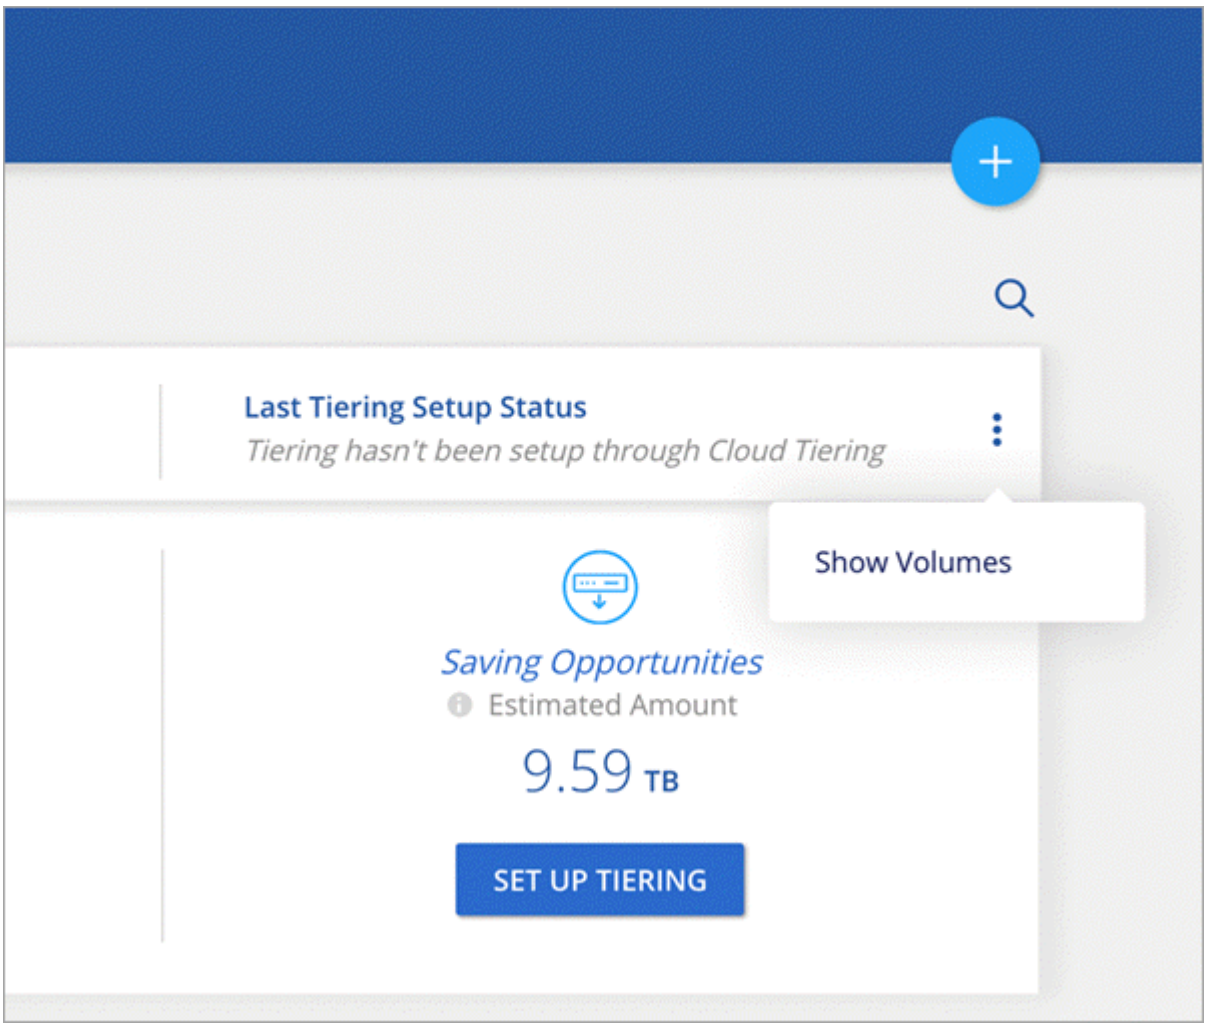

Multi-tenancy using Cloud Central accounts

Cloud Tiering now supports multi-tenancy through Cloud Central accounts. Accounts enable multiple users to manage the same clusters in an account.

Learn more about accounts.

18 May 2020

This update includes the following new features and enhancements.

Cooling days

Cloud Tiering now enables you to adjust the cooling period that's associated with a volume tiering policy. The cooling period is the number of days that user data in a volume must remain inactive before it's considered "cold" and moved to object storage. Adjusting the cooling days enables you to optimize tiering for your environment.

Learn how to adjust the cooling period.

Rediscovery of clusters

We've made improvements to the Cloud Tiering service that requires you to deploy new Service Connectors and rediscover your clusters.

The benefits include improved performance, the ability to release new features faster, improved error handling, an improved discovery process, and more.

Cloud Tiering will prompt you to rediscover your clusters when you log in. If you need help, contact us by using the in-product chat.

1 Apr 2020

We're pleased to announce the following new features and enhancements.

Simplified cluster setup

We enhanced how you set up data tiering on a new cluster. The simplified wizard walks you through three steps: discover the cluster, set up the cluster for tiering, and then choose the volumes that you want to tier. The following pages provide more details:

- Tiering data to AWS S3
- Tiering data to Azure Blob storage
- Tiering data to Google Cloud storage

Ability to add or select an existing bucket/container

When you set up data tiering on a cluster, you now have the choice to add a new bucket/container or to select an existing bucket/container. Prior to this change, Cloud Tiering automatically created it for you.

Support for additional S3 storage classes

Cloud Tiering now supports two additional S3 storage classes: *Zone-IA* and *Intelligent*.

When you set up data tiering, Cloud Tiering can apply a lifecycle rule so the tiered data transitions from the *Standard* storage class to another storage class after 30 days. Learn more.

25 Dec 2019

This update includes the following new features and enhancements.

Changing a volume's tiering policy

You can now change a volume's tiering policy after you set the initial tiering policy. For example, you can easily change from *Cold snapshots* to *Cold user data*.

Learn how to change a volume's tiering policy.

Licensing enhancements

- Support for ONTAP FabricPool licenses

  Cloud Tiering now supports the FabricPool licenses that you add to an ONTAP cluster. If a license is available on the cluster, then Cloud Tiering displays the license in the **Licensing** page and enables you to tier data to the cloud based on that license.

- Support for pay-as-you-go from the Azure Marketplace

  You can now pay-as-you-go when tiering cold data to Azure Blob storage.

- Support for pay-as-you-go from the GCP Marketplace

  You can now pay-as-you-go when tiering cold data to Google Cloud storage.

- Support for a combination of PAYGO and BYOL licensing

  You can combine a FabricPool license with a pay-as-you-go subscription to pay for data that you tier to the cloud. If your FabricPool license expires or if you tier more data than allowed by the license, then data tiering is never interrupted—it continues through your PAYGO subscription.

For more details:

- Learn how licensing works
- Learn how to set up pay-as-you-go subscriptions and BYOL

3 Nov 2019

This update includes the following new features and enhancements.

Support for Google Cloud

You can now tier inactive data from your ONTAP clusters to Google Cloud Storage.

Learn how to get started

Support for additional tiering policies

You can now select the following tiering policies when setting up tiering:

- All user data (All)
- All DP user data (Backup)

Learn about these tiering policies.

Tiering policy per volume

You can now choose a different tiering policy for each volume when you set up tiering.

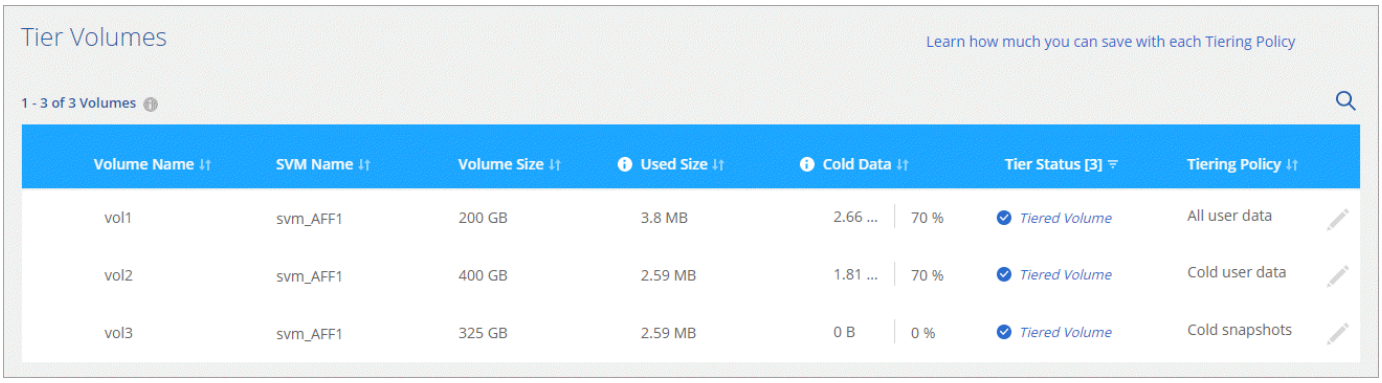

8 Sept 2019

This Cloud Tiering update includes the following new features and enhancements.

Tiering setup status

Cloud Tiering now shows the status of tiering setup for each cluster. For example, the status might indicate that Cloud Tiering is setting up the object store, or that it successfully set up 145 volumes for data tiering. The status also identifies if any failures occurred during setup.

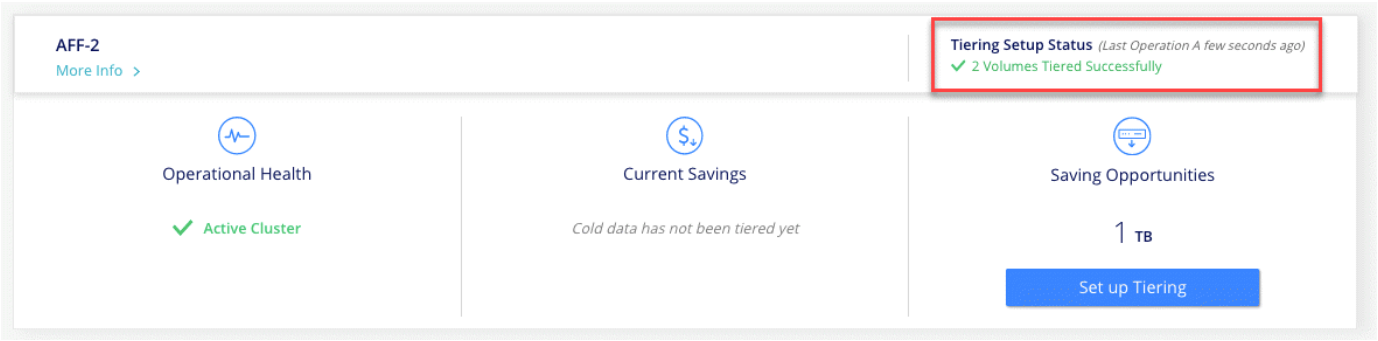

Integration with Cloud Manager AWS subscriptions

If you use NetApp Cloud Manager and you've already subscribed through its new AWS Marketplace offering, then you're automatically subscribed to Cloud Tiering, as well. You'll see an active subscription in Cloud Tiering in the **Licensing** tab. You won't need to subscribe.

If you've already subscribed through Cloud Tiering, then this change has no impact to you. You're all set.

7 Aug 2019

This update includes the following new features and enhancements.

On-premises Service Connector for data tiering to S3

You can now install the Service Connector on an on-premises Linux host when you tier cold data to AWS S3. Learn more.

Auto discovery of clusters

If you discovered ONTAP clusters through NetApp Cloud Manager, the clusters are automatically added to Cloud Tiering if they support data tiering.

Object storage connectivity check

If Cloud Tiering finds a connectivity problem with the object storage bucket, the tiering health indicator in the dashboard provides details about the problem.

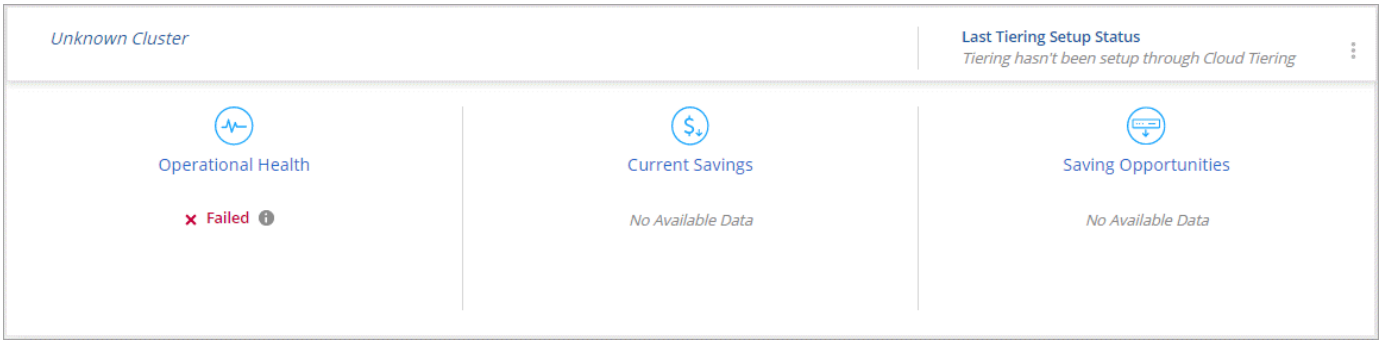

4 July 2019

Cloud Tiering was updated to fix a few bugs.

10 June 2019

This update includes the following new features and enhancements.

Cloud Tiering is now generally available

The Controlled Availability release of Cloud Tiering has completed—Cloud Tiering is now available for customer use from NetApp Cloud Central. A 30-day free trial is available for both AWS and Azure. It starts when you set up tiering to your first cluster.

Pay-as-you-go from the AWS Marketplace

After your free trial starts, subscribe to the Cloud Tiering service to ensure that there's no disruption of service after the trial ends. When it ends, you'll be charged hourly according to the amount of data that you tier.

Learn how to subscribe from the AWS Marketplace.

We're planning to add Cloud Tiering to the Azure Marketplace as soon Azure supports SaaS pricing.

Support for FlexGroup volumes

You can now tier inactive data from FlexGroup volumes to object storage, starting with ONTAP 9.5. Setup works the same as any other volume.

5 May 2019

This update includes the following new features and enhancements.

Support for Microsoft Azure

You can now tier inactive data from your ONTAP clusters to Azure Blob storage.

- Learn how to tier inactive data to Azure
- Review support for Azure Blob access tiers and Azure regions

Ability to choose an IPspace for connections to object storage

When you set up tiering for an ONTAP cluster, you now need to select the IPspace that ONTAP should use to connect to object storage. Selecting the correct IPspace ensures that Cloud Tiering can set up a connection from ONTAP to your cloud provider's object storage.

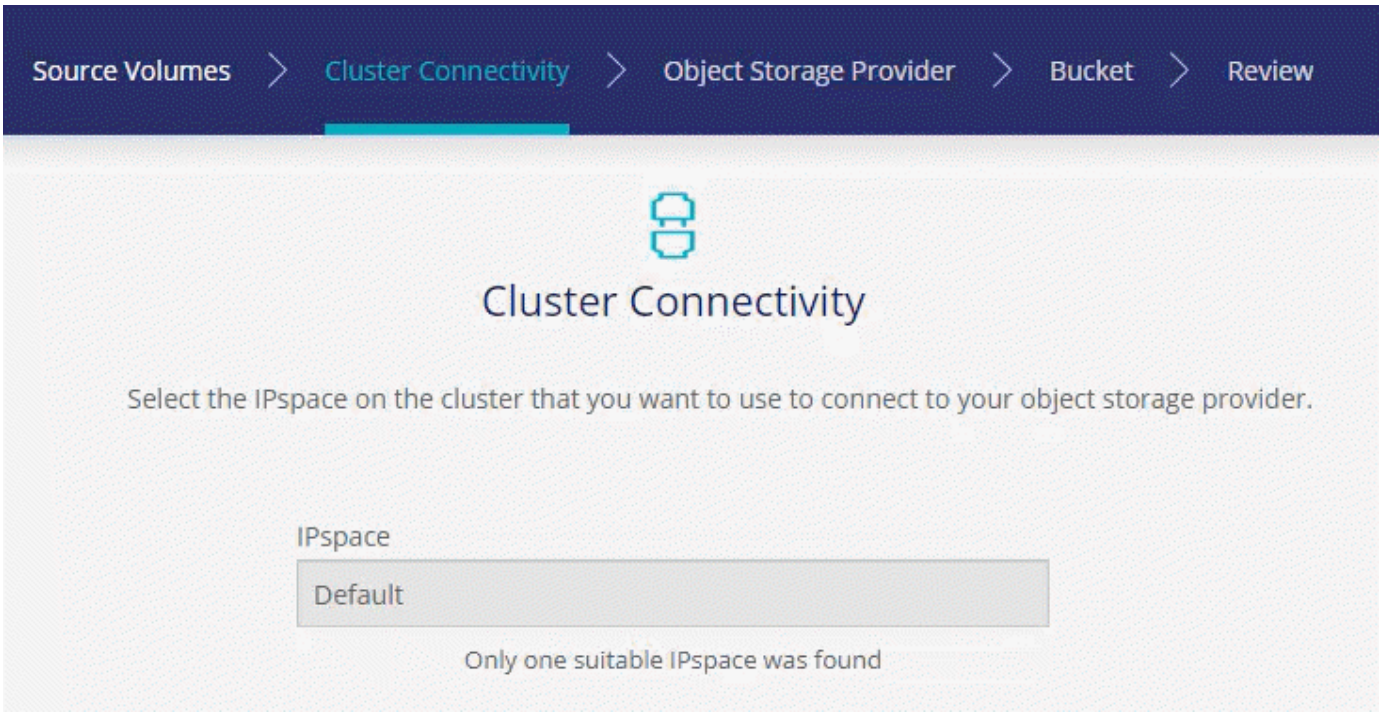

To understand the requirements for the IPspace and the associated intercluster LIFs, refer to ONTAP cluster requirements:

- Preparing to tier inactive data to AWS S3
- Preparing to tier inactive data to Azure Blob storage

7 Apr 2019

This update includes the following new features and enhancements.

- Support for FAS systems with all-SSD aggregates
- Support for additional versions of ONTAP
- Ability to choose the type of cold data that you want to tier
- Ability to choose an S3 storage class

Support for FAS systems with all-SSD aggregates

In addition to AFF systems, Cloud Tiering now supports FAS systems that have one or more all-SSD aggregates.

Support for additional versions of ONTAP

Cloud Tiering now supports ONTAP 9.2 and 9.3. This is in addition to supporting ONTAP 9.4 and later.

Ability to choose the type of cold data that you want to tier

For ONTAP 9.4 and later, you can now choose the type of cold data that you want to tier. You can tier *all cold data* or just *Snapshot copies*.

If you have an ONTAP 9.2 or 9.3 system, you can tier Snapshot copies only.

The option is available after you click **Set up Tiering**.

- Learn more about the cold data that you can tier
- Learn how to tier data from your first cluster
- Learn how to tier data from additional volumes

Ability to choose an S3 storage class

When you set up data tiering on a cluster for the first time, you can now choose to tier cold data to the S3 *Standard* storage class or to the *Standard-Infrequent Access* storage class. Learn about these S3 storage classes.

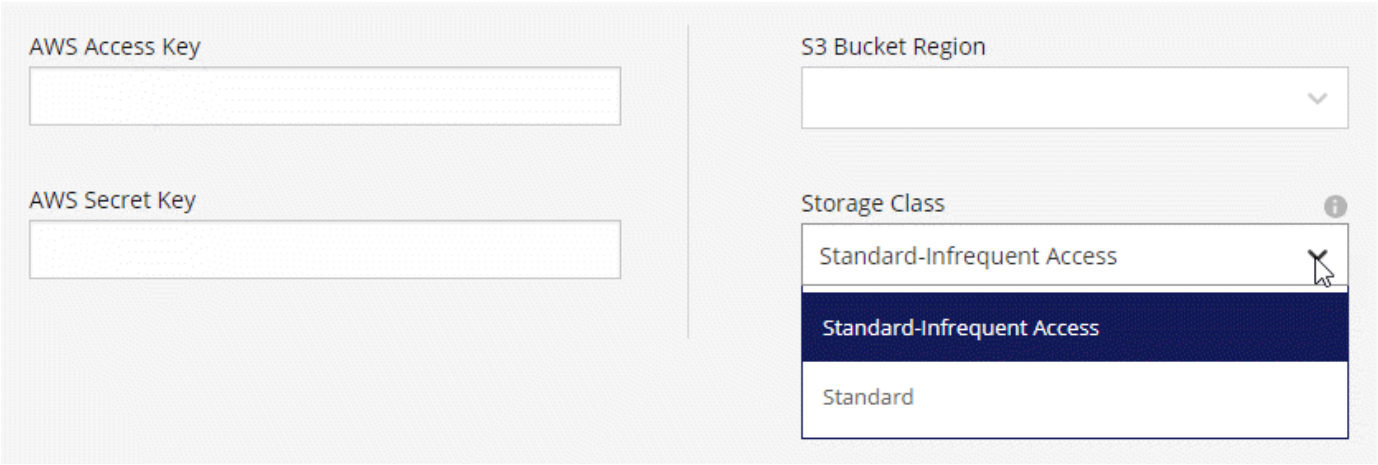

Learn how to set up data tiering on a cluster.

Copyright Information

Copyright © 2021 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system-without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP “AS IS” AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at http://www.netapp.com/TM are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.