Tiering inactive data from your first cluster to AWS S3

Cloud Tiering

Ben Cammett, Aksel Davis
October 07, 2020

This PDF was generated from https://docs.netapp.com/us-en/cloud-tiering/task_tiering.html on March 06, 2021. Always check docs.netapp.com for the latest.

Table of Contents

Tiering inactive data from your first cluster to AWS S3 1

Tiering inactive data from your first cluster to AWS S3

After you prepare your AWS environment, just log in to Cloud Tiering and start tiering inactive data from your first cluster.

What you'll need

- To discover the cluster, you'll need the following:
  - The cluster management IP address.
  - The user name and password of an ONTAP account that has administrator-level privileges.

    The Service Connector uses this account to send API operations to the ONTAP cluster.
- To deploy the Service Connector in AWS, you'll need the following:
  - The AWS region, VPC, and subnet in which the Service Connector will reside.
  - An AWS access key for an IAM user who has the required permissions.
- To set up tiering to S3, you'll need an AWS access key for an IAM user who has the required S3 permissions.

If you haven't met these requirements, see Preparing your environment.

Steps

1. Log in to NetApp Cloud Central.
2. Select the Cloud Tiering service.
3. Click **Let's Start, Discover Your First Cluster**.
4. Complete the steps on the **Discover Cluster** page:
   a. Enter the cluster management IP address and the user name and password of an account that has administrator-level privileges.

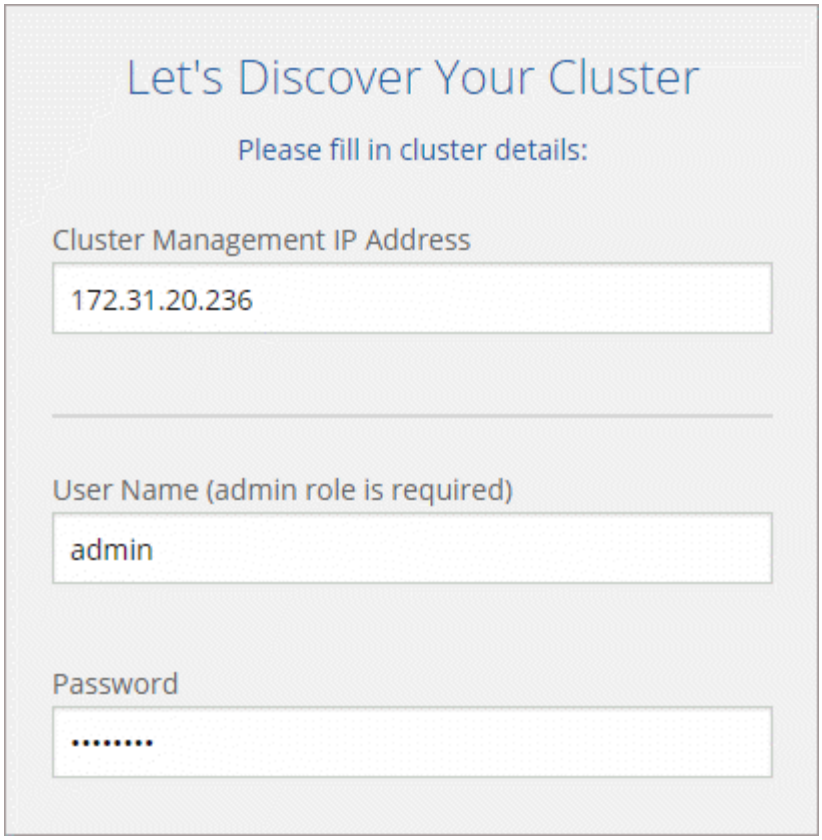

b. Click **Discover Cluster**.

   If you already have an existing Service Connector, then Cloud Tiering automatically attempts to use that Service Connector. Cloud Tiering moves on to the next step, if any existing Service Connector has connectivity to the cluster.

c. If Cloud Tiering prompts you to create a Service Connector, click **Create your first Service Connector**.

   If you have an existing Service Connector and Cloud Tiering can't use it, then you'll need to do one of two things:

   - Click **Add Service Connector** to create a new Service Connector.
   - Check the status and connectivity of any existing Service Connectors and then try to discover the cluster again.

d. If you clicked the button to create a Service Connector, follow the prompts to deploy it in AWS:

   - **Select Provider**: Select **Amazon Web Services** as the target location for the Service Connector.
   - **AWS Credentials**: Enter the AWS access key ID and secret key for an IAM user that has the required permissions to deploy the Service Connector.
   - **Location**: Select the AWS region, VPC, and subnet for the Service Connector EC2 instance.

     Remember, the Service Connector must have a constant connection to the ONTAP cluster and a constant internet connection to the Cloud Tiering service.

   - **Network**: Select a key pair to use for the EC2 instance, choose whether to assign a public IP, and specify an HTTP proxy, if one is required for outbound connectivity.
   - **Security Group**: Select **Create a new security group** so Cloud Tiering can create the security

group, or select your own. Then click **Go**.

The security group that Cloud Tiering creates has no inbound connectivity and open outbound connectivity.

Leave the page open until the deployment is complete.

e. Back on the Discover Cluster page, select the Service Connector that you just created.

5. Complete the steps on the **Tiering Setup** page:

   a. **S3 Bucket**: Add a new S3 bucket or select an existing S3 bucket that starts with the prefix *fabric-pool* and click **Continue**.

      The *fabric-pool* prefix is required because the IAM policy for the Service Connector enables the instance to perform S3 actions on buckets named with that exact prefix.

      For example, you could name the S3 bucket fabric-pool-AFF1, where AFF1 is the name of the cluster.

   b. **Storage Class**: Select the S3 storage class that you want to transition the data to after 30 days and click **Continue**.

      If you choose Standard, then the data remains in that storage class.

   c. **Credentials**: Enter the access key ID and secret key for an IAM user who has the required S3 permissions.

      The IAM user must be in the same AWS account as the bucket that you selected or created on the **S3 Bucket** page.

   d. **Cluster Network**: Select the IPspace that ONTAP should use to connect to object storage and click **Continue**.

      Selecting the correct IPspace ensures that Cloud Tiering can set up a connection from ONTAP to your cloud provider's object storage.

      If you haven't reviewed requirements for the IPspace and the associated intercluster LIFs, see ONTAP cluster requirements.

6. Click **Continue** to select the volumes that you want to tier.

7. For each volume, click the icon, select a tiering policy, optionally adjust the cooling days, and click **Apply**.

   Learn about volume tiering policies and cooling days.

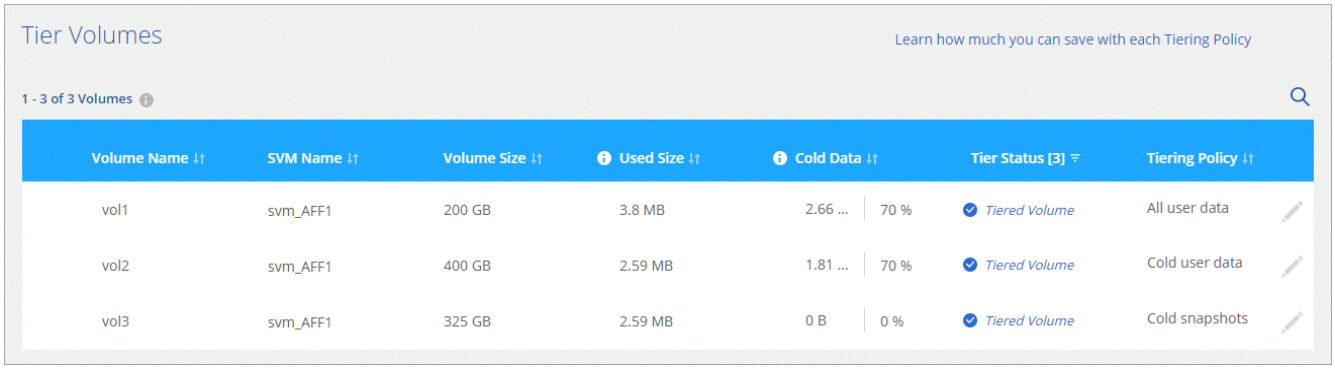

8. When you’re done, click **Close**.

Result

You’ve successfully set up data tiering from volumes on the cluster to S3 object storage.

What’s next?

Be sure to subscribe to the Cloud Tiering service from the AWS Marketplace.

You can also add additional clusters or review information about the active and inactive data on the cluster. For details, see Managing data tiering from your clusters.

Copyright Information

Copyright © 2021 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system-without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP “AS IS” AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at http://www.netapp.com/TM are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.